OpenClawBrain v12.2.6+: Shadow Routing with QTsim Confidence Mixing and Unified Policy Learning

Jonathan Gu

March 2, 2026

Version: v12.2.6+

1 Abstract

OpenClawBrain is a retrieval-routing system for long-lived agents. Query-time retrieval stays local and deterministic, then offline learning updates routing policy from multiple supervision sources. The current system uses two route signals: a structural prior (graph_prior) and a query-time compatibility signal (QTsim). A confidence mixer combines those signals using uncertainty features, including entropy and margin, to decide how much weight each source gets per decision. Learning is performed through teacher distillation plus policy-gradient updates under a reward-source hierarchy. Human corrections have highest authority, then self-learning outcomes, then harvested signals, then high-volume teacher labels. This release also starts RL-native maintenance of graph structure: Phase 2a is shipped and follow-on phases cover connect/split/merge actions under policy control.

2 1. System Overview

OpenClawBrain is designed around a strict split:

- **Hot path:** local embedding retrieval and bounded traversal only.
- **Cold path:** asynchronous learning and structure updates.

This split keeps production latency stable while allowing large-volume policy updates in the background.

1.1 Routing Inputs

At each route decision over candidate targets $\mathcal{A}(s)$ from state s:

- $\text{graph_prior}(a \mid s)$ captures persistent graph structure.
- $\text{QTsim}(a \mid q, s)$ captures query-conditioned fit.

1.2 Confidence Mixing

The router computes a query-local mix coefficient $\lambda \in [0, 1]$ from confidence features. Two required features are:

- **Entropy** of QTsim over candidates (high entropy means low confidence).
- **Margin** between top-1 and top-2 QTsim scores (low margin means ambiguity).

A simple form is:

$$\lambda = \sigma\big(\alpha_0 + \alpha_1 \,\mathrm{margin}_{\mathrm{QTsim}} - \alpha_2 \, H(\mathrm{QTsim})\big)$$

$$\pi(a \mid q, s) = \lambda \,\mathrm{QTsim}(a \mid q, s) + (1 - \lambda)\,\mathrm{graph_prior}(a \mid s)$$

where σ is a logistic gate and H is entropy. When QTsim is sharp and separated, λ increases. When QTsim is uncertain, the policy falls back toward graph prior.

3 2. Learning Loop: Distillation + Policy Gradient

The learning loop combines distillation-style supervision with policy-gradient refinement.

2.1 Teacher Distillation

Offline replay samples route decisions and asks a teacher policy for preferred actions (or distributions). Distillation improves route quality quickly and provides dense supervision for states with sparse human labels.

2.2 Policy-Gradient Update

Given route trajectory $\tau = (s_0, a_0, \ldots, s_T, a_T)$ and scalar return $R(\tau)$:

$$\nabla J(\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[\sum_{t=0}^{T} (R(\tau) - b_t) \nabla_\theta \log \pi_\theta(a_t \mid s_t) \right]$$

where b_t is a baseline. Distillation targets and reward-weighted gradients are both applied; the exact blend is implementation-defined per run mode.

2.3 Reward-Source Hierarchy

The system resolves conflicting supervision by ordered authority:

1. Human correction/approval (highest authority)
2. Self-learning outcomes from deployed agents
3. Harvested weak signals
4. Teacher labels from async replay (lowest per-sample authority, highest volume)

This hierarchy constrains updates so high-authority signals can override lower-level noise.

4 3. RL-Native Maintain (Structure Learning)

OpenClawBrain treats structure maintenance as a control problem over graph actions.

3.1 Phase 2a (Shipped)

Phase 2a ships RL-native maintenance hooks that integrate policy outcomes with structure scoring and safe-apply gates. In practice, this means structure candidates are generated, scored with policy-derived signals, and applied only when guardrails pass.

3.2 Roadmap: connect/split/merge

Planned phases extend action space with explicit graph edits:

- **Connect:** create/reweight candidate edges when repeated successful transitions lack direct structure.
- **Split:** divide overloaded nodes when policy confidence or error patterns indicate topic collapse.
- **Merge:** combine nodes with persistent redundancy and equivalent downstream behavior.

Each action class is expected to use conservative thresholds, rollback support, and offline shadow evaluation before broad enablement.

5 4. Evaluation Plan (Aligned to `docs/evaluation-plan.md`)

Evaluation is organized around reproducible, decision-driving checks.

4.1 Primary Questions

- Does QTsim confidence mixing outperform fixed-weight or single-signal routing?
- Does distillation+policy-gradient outperform either component alone?
- Do RL-native maintenance actions improve long-run retrieval quality without instability?

4.2 Core Metrics

- Task success / correctness rate (dataset-dependent)
- Retrieval precision/recall over required context nodes
- Token/context budget per task and latency impact
- Route stability (variance, oscillation, regression rate)
- Maintenance safety (rollback rate, bad-edit incidence)

4.3 Ablations

Minimum ablations to report:

- graph_prior only vs QTsim only vs mixed policy
- fixed mix vs entropy/margin confidence gating
- distillation only vs policy gradient only vs combined
- no-maintain vs connect-only vs connect+split+merge (as phases ship)

4.4 Reporting Template (Do Not Invent Numbers)

All unreproduced results must be marked `TBD`. Example table:

Setting	Success (%)	Avg Context Tokens	Latency (ms)
graph_prior only	TBD	TBD	TBD
QTsim only	TBD	TBD	TBD
mixed (entropy+margin gate)	TBD	TBD	TBD

Table 1: Populate from reproducible runs only. Replace each `TBD` with run ID, date, and command reference.

When filling `TBD` values:

- record exact commit SHA,
- include command lines and seed settings,
- link run artifacts/logs,
- report confidence intervals where supported.

6 5. Operational Notes

OpenClawBrain now supports local embeddings out-of-the-box for operator-first deployment. This enables deterministic bootstrap and private/offline-friendly operation while keeping optional teacher workflows in asynchronous jobs.

7 6. Conclusion

The v12.2.6+ system is defined by four commitments: local hot-path retrieval, QTsim+graph-prior confidence mixing, distillation plus policy-gradient learning, and RL-native structure maintenance. Performance claims should be published only from reproducible evaluation runs with explicit provenance.

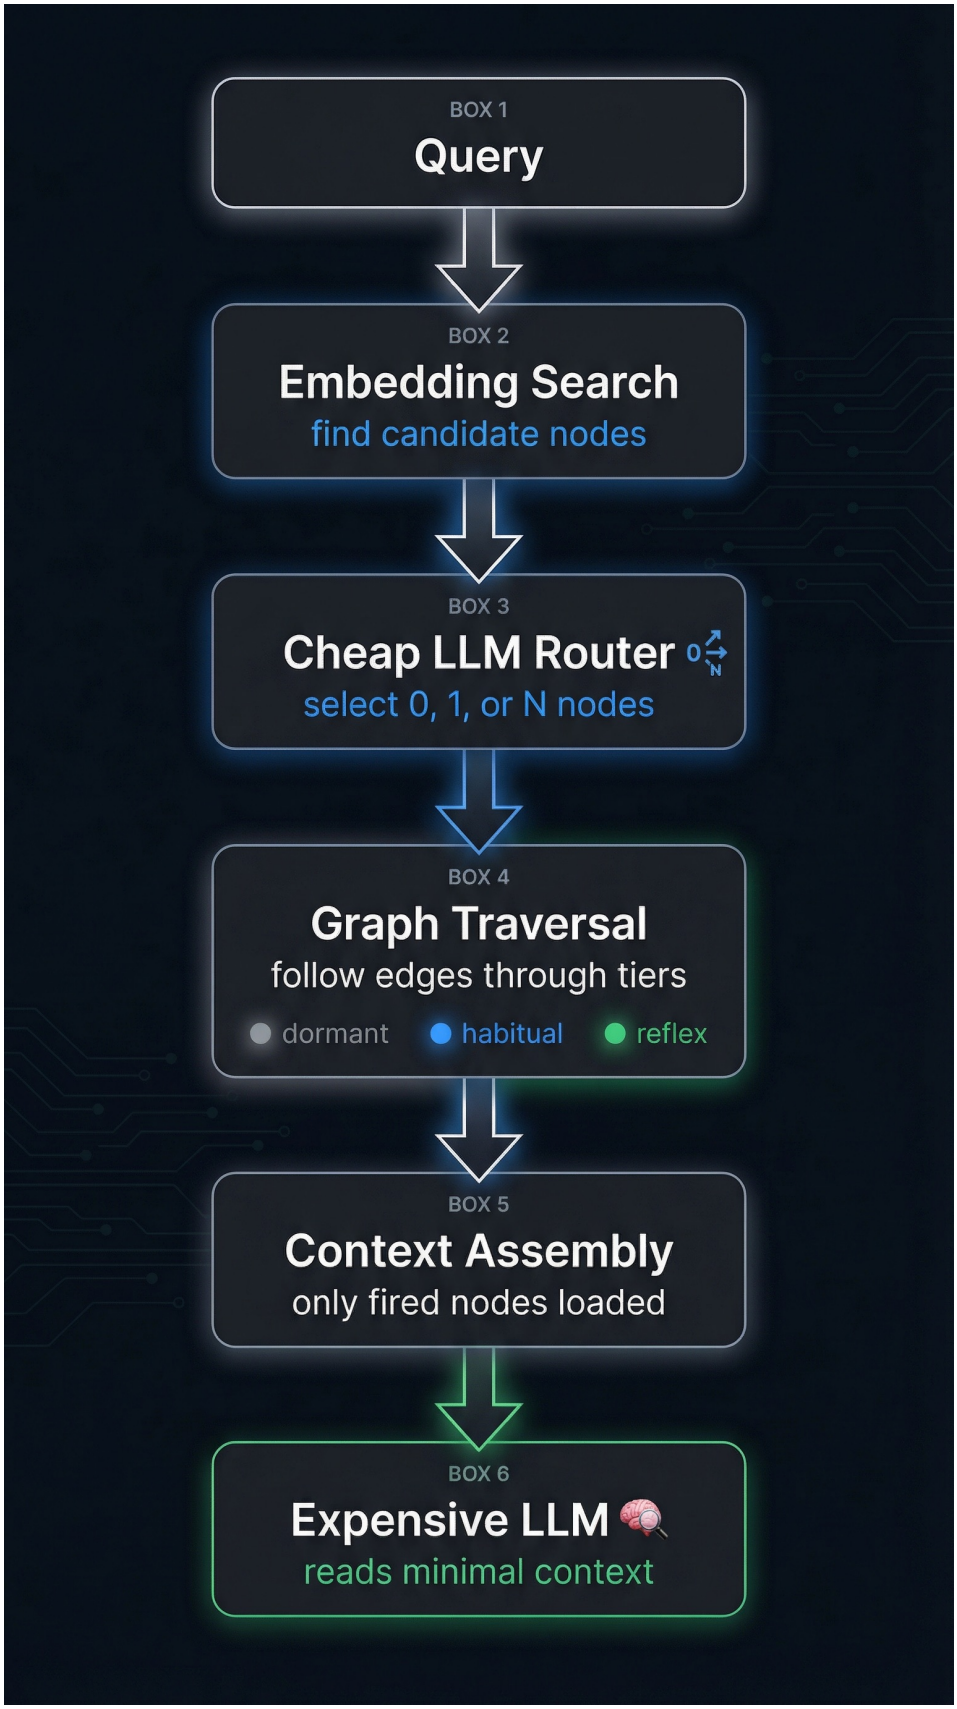

Figure 1: System diagram: latency-critical hot path stays local and bounded, while shadow replay and policy updates run asynchronously in the cold path.

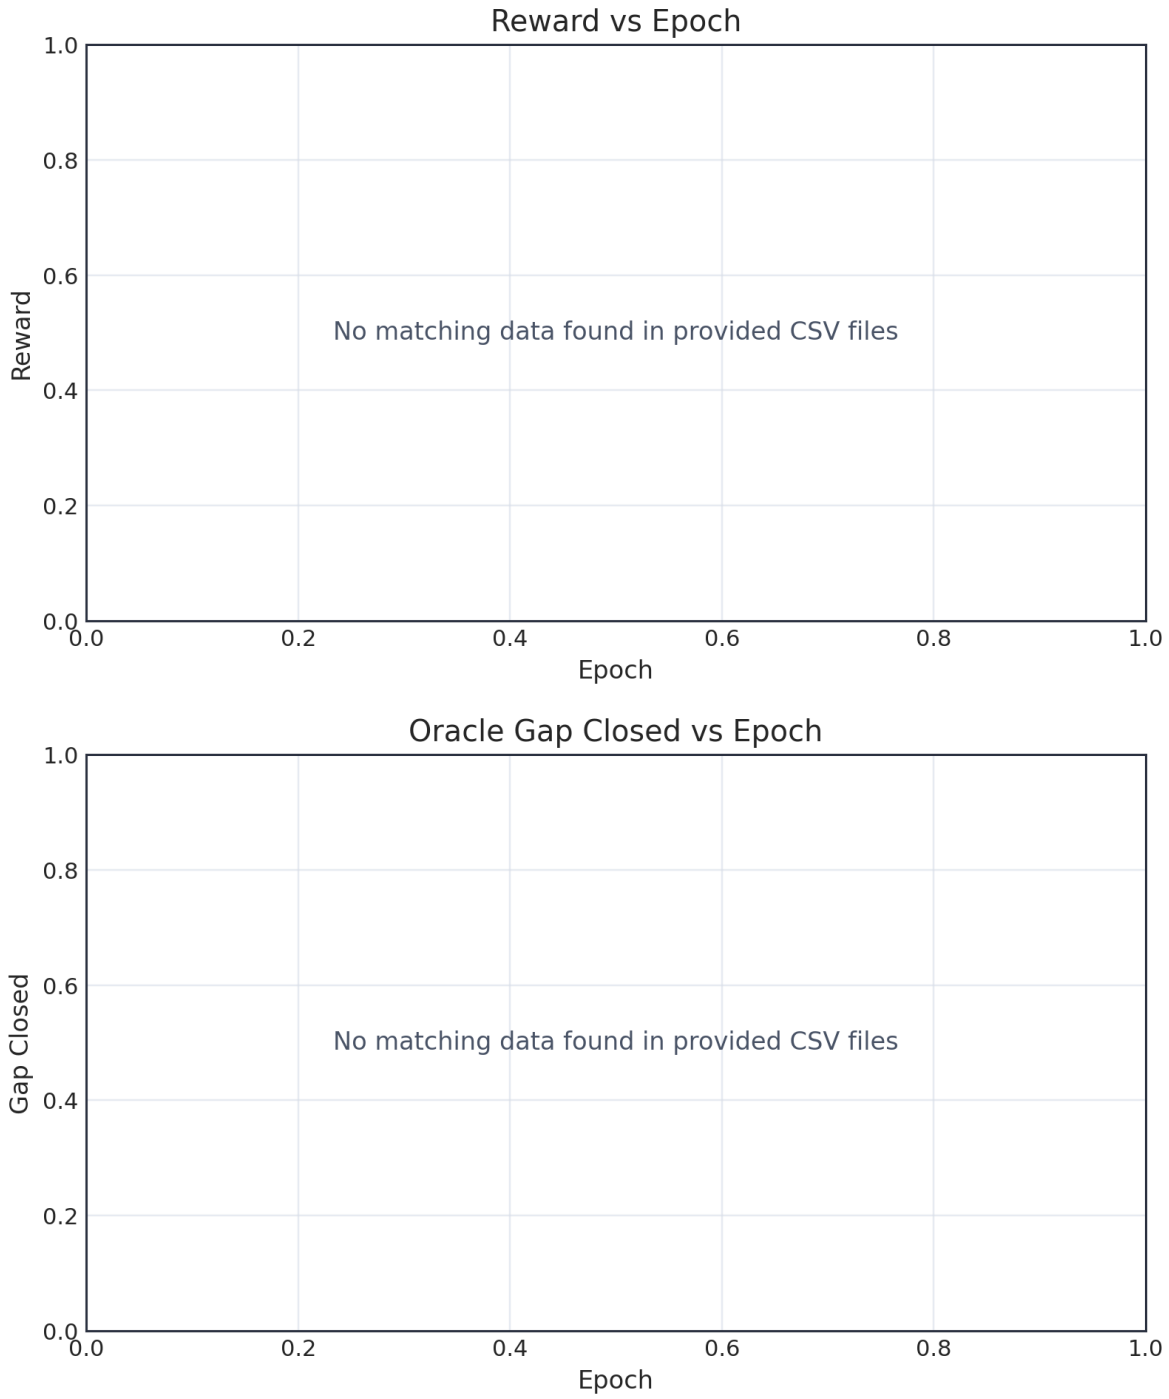

Figure 2: Expert-regions simulation learning behavior: reward trajectory and fraction of oracle gap closed across training epochs.

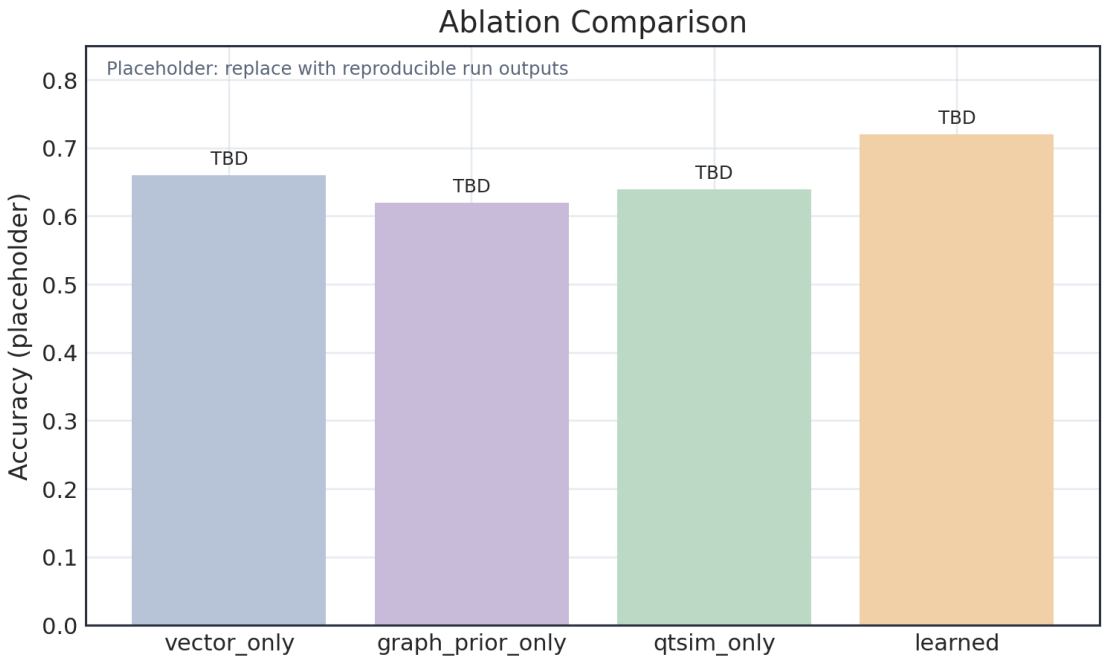

Figure 3: Ablation comparison across vector_only, graph_prior_only, qtsim_only, and learned routing. Placeholder bars are used until reproducible ablation CSV outputs are published.